

# Group: 22

# Project: Text2SQL Demo

Jinal Vyas

Nishtha Kuberbhai Chaudhary

Omkar Shelar

Rebecca Pires Dos Santos

Vrishir Bhaskar Iyer

# Presentation Outline

- Problem Definition
- System Architecture & Algorithms
- Datasets (Description, Size and Preprocessing)
- Evaluation (Metrics, Experiments, Preliminary Findings)
- UI Interfaces or Data Visualization (Preliminary Designs)
- Project Plan: Tasks, Deadlines, Division of Work

# Problem Definition

Existing large language models like ChatGPT and Gemini, while powerful, are not optimized for schema-aware Text-to-SQL translation and often produce semantically invalid queries. They also pose issues of cost, privacy, and limited customization.

**Solution:** To address this, we develop a custom Text2SQL pipeline and fine-tune an LLM for domain-specific SQL generation. A traditional ML model is used as a baseline to evaluate and highlight the performance gains achieved through LLM fine-tuning.

# Dataset Description

- Spider 1.0
  - <https://yale-lily.github.io/spider>
  - Complex and cross-domain text-to-SQL dataset.
  - 10,181 questions and 5,693 SQL queries on 200 databases.
  - The queries are unique and complex.
  - Multiple tables which cover 138 different domains.
- BIRD
  - <https://bird-bench.github.io/>
  - Targets real-world conditions: larger, messier databases, dirty or ambiguous values, need for external knowledge, and SQL efficiency
  - BIRD stresses robustness, execution, and scalability beyond Spider 1.0.

# System Architecture for Data-Segmentation Field

- We followed two approaches to better our previously fine-tuned LLM model
  - Natural Language Query based clustering
  - SQL Query based clustering
- Both were done against train data samples (disjoint from test data samples) to allow for schema generalization
- We feed the DB schema as well as the question for CodeLlama to better understand the DB structure
- Reasons to cluster:
  - SPIDER has diverse SQL patterns so it helps to create “specialists” models to help with similar queries.

# Fine-Tuning based on Natural Language

## 1. Clustering:

- Extract question from prompt
- Embed using MiniLM
- Run K-Means ( $k = 4$ )
- Create clusters and centroids

## 2. Cluster-Specific Fine-Tuning:

- Fine-tune one LoRA per cluster

## 3. Inference time Router: Selecting the Right cluster/LoRA

Question → nearest centroid → pick LoRA + model → SQL

## 4. Use the correct LoRA adaptor with CodeLlama to get SQL

```
((base) [jjvyas1@sol-login01:~/codellama_finetune/processed_spider/clusters]$ python cluster_display.py

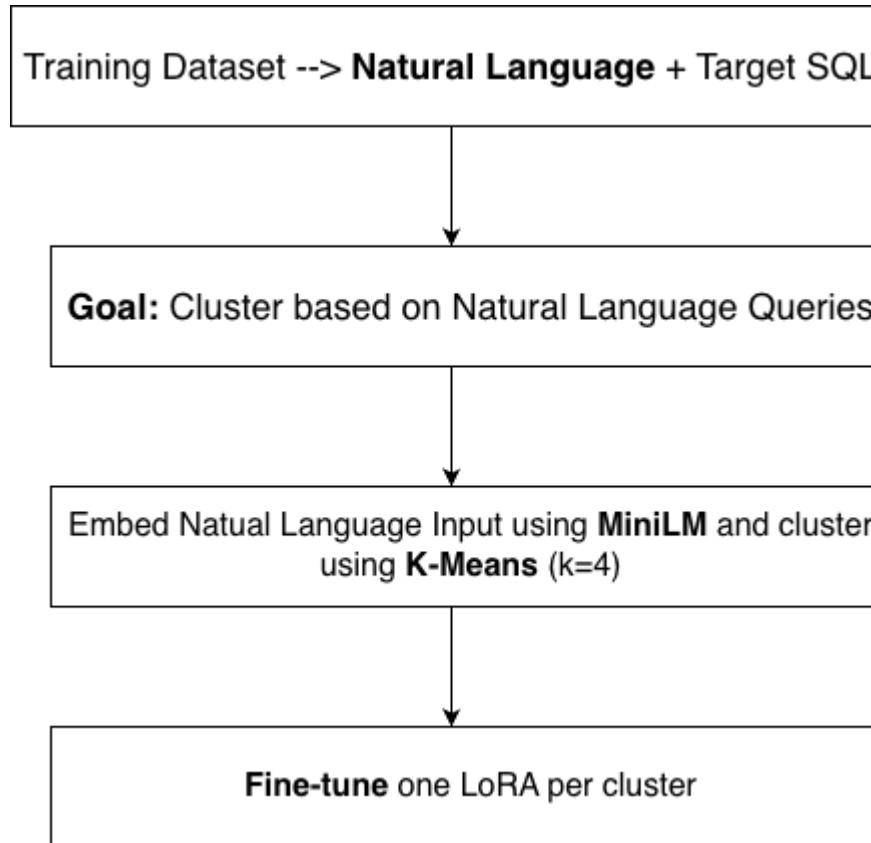
== train_cluster_0.jsonl ==
1. How many heads of the departments are older than [VAL_1] ?
2. What is the average number of employees of the departments whose rank is between [VAL_1] and [VAL_2]?
3. What are the distinct creation years of the departments managed by a secretary born in state [VAL_1]?
4. What are the names of the states where at least [VAL_1] heads were born?
5. Show the name and number of employees for the departments managed by heads whose temporary acting value is [VAL_1]?

== train_cluster_1.jsonl ==
1. List the name, born state and age of the heads of departments ordered by age.
2. List the creation year, name and budget of each department.
3. What are the maximum and minimum budget of the departments?
4. Count the number of farms.
5. What are the maximum and minimum number of cows across all farms.

== train_cluster_2.jsonl ==
1. What are the names of the heads who are born outside the California state?
2. How many acting statuses are there?
3. How many departments are led by heads who are not mentioned?
4. What are the distinct ages of the heads who are acting?
5. How many farms are there?

== train_cluster_3.jsonl ==
1. In which year were most departments established?
2. Which course has most number of registered students?
3. What is the name of the course with the most registered students?
4. What is id of students who registered some courses but the least number of courses in these students?
5. What are the ids of the students who registered for some courses but had the least number of courses for all students?
```

# Fine-Tuning based on Natural Language



# Fine-Tuning based on Target SQL

Training Dataset --> Natural Language + **Target SQL**

**Goal: Clustering** Target SQLs to remove outliers  
and better fine-tune LLMs

SQL -> Syntax Tree (`sqlglot`) -> Simplify it to remove the  
table names, attribute names, literals (similar to AST)

Syntax Tree to a string (**TRC-like syntax**)

Convert to word-embeddings and cluster it

- select T2.first\_name , T2.last\_name from candidates AS T1  
join people AS T2 on T1.candidate\_id = T2.person\_id

**AND ( IDENT\_1 . IDENT\_4 EQ IDENT\_2 . IDENT\_3 ) PROJ  
PROJ\_COL PROJ\_COL**

- select T2.Name , count(\*) from debate\_people AS T1 join  
people AS T2 on T1.Affirmative = T2.People\_ID group by  
T2.Name

**AND ( IDENT\_4 . IDENT\_6 EQ IDENT\_2 . IDENT\_7 )  
GROUPBY IDENT\_2 . IDENT\_1 PROJ PROJ\_COL AGG\_\***

# Fine-Tuning based on Target SQL (contd.)

- Algorithms used for clustering:
  - HDBScan – Too much concentration of clusters

```
{  
    "num_clusters_excluding_noise": 3,  
    "num_points_total": 6492,  
    "num_outliers": 73,  
    "points_per_cluster": {  
        "-1": 73,  
        "0": 10,  
        "1": 6401,  
        "2": 8  
    }  
}
```

# Fine-Tuning based on Target SQL (contd.)

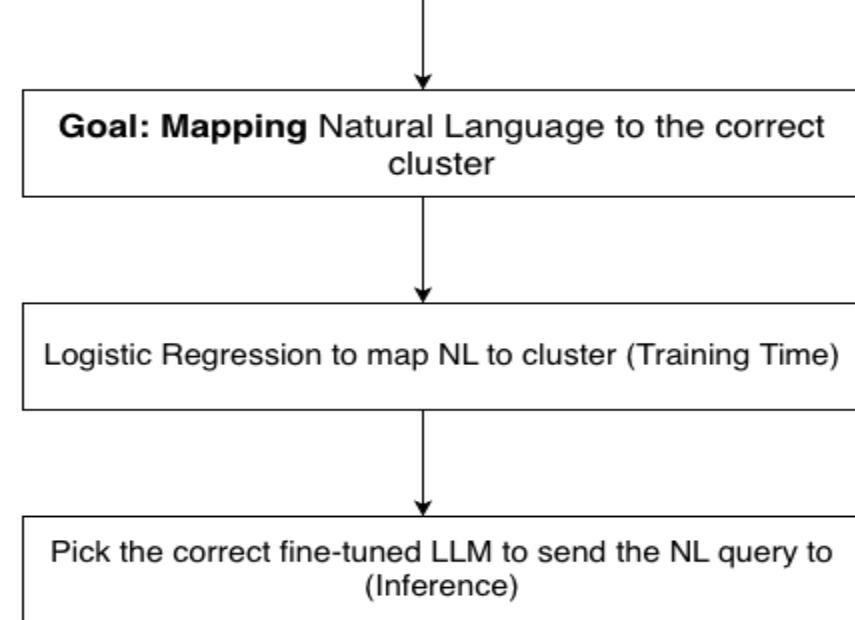
- Algorithms used for clustering:
  - Hierarchical (agglomerative) Clustering

```
{  
    "num_clusters_excluding_noise": 3,  
    "num_points_total": 6492,  
    "points_per_cluster": {  
        "0": 1537,           # Simple projections / aggregations  
        "1": 3223,           # Joins and complex selection  
        "2": 1732            # Grouping and Aggregations  
    }  
}
```

# Fine-Tuning based on Target SQL (contd.)

Training Dataset --> **Natural Language + Target SQL**

Validation accuracy: 0.7662  
Training:Validation Split => 80:20



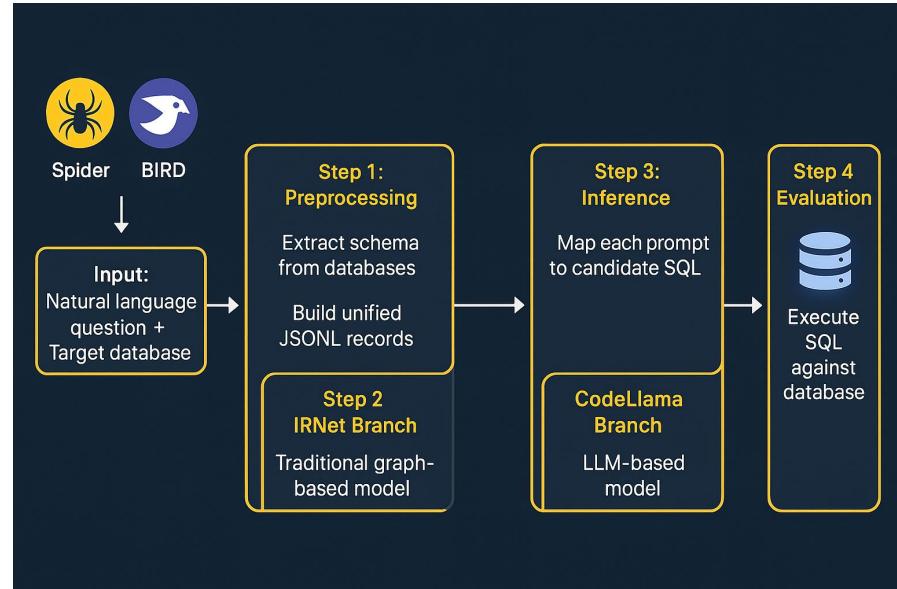
	precision	recall	f1-score
0	0.71	0.67	0.69
1	0.82	0.87	0.84
2	0.71	0.66	0.68

accuracy			0.77
macro avg	0.74	0.73	0.74
weighted avg	0.76	0.77	0.76

# System Architecture for LLM-Spider and BIRD

## Phase 1 – CodeLlama with Spider 1.0

- **Goal:** Establish a baseline Text-to-SQL performance on a standard academic benchmark.
- **Dataset:** Spider 1.0 – 200 databases, 138 domains, complex cross-domain SQL queries.
- **Pipeline:**
  - Preprocess Spider → serialize schema + question into a unified prompt.
  - Fine-tune **CodeLlama-7B** using LoRA on train split.
  - Generate SQL for dev split; evaluate with **EX**.
- **Outcome:**
  - Demonstrates how well CodeLlama learns structured SQL on a *clean, curated* benchmark.
  - Builds our reference point for later comparison with BIRD.



# Phase 2 – CodeLlama with BIRD

**Goal:** Test CodeLlama on a **large, realistic** benchmark.

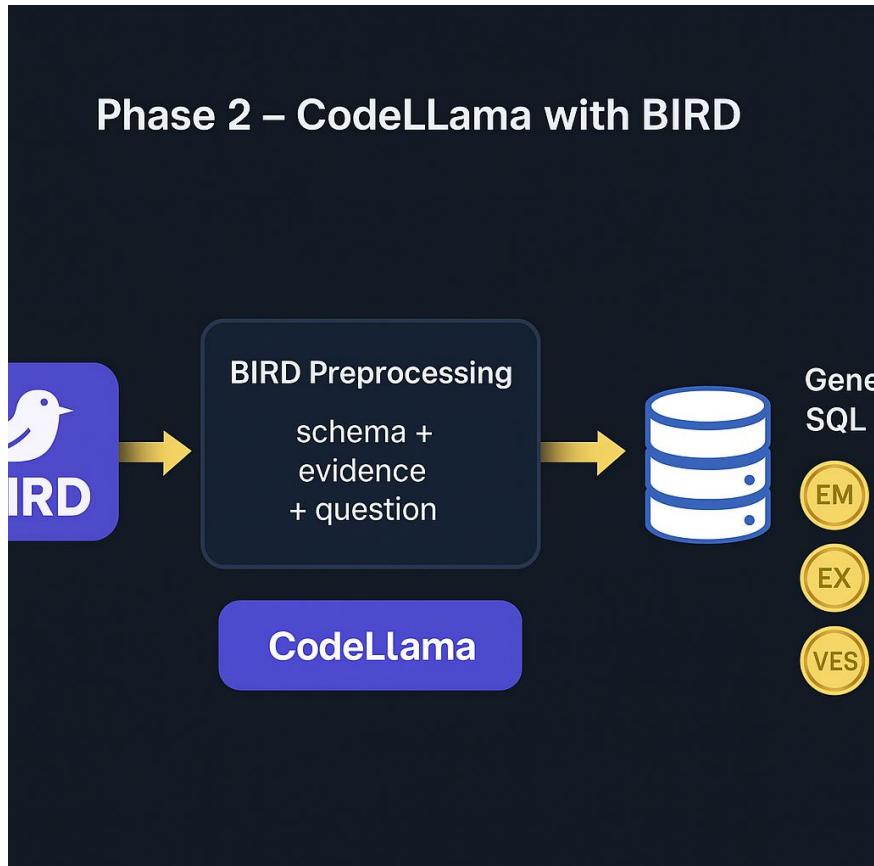
**Dataset:** BIRD – 12,751 question–SQL pairs, 95 large databases, 33.4 GB across 37+ professional domains (finance, healthcare, education, etc.).

**Key Changes vs Spider:**

- Much **larger schemas and table counts** → heavier prompts.
- Additional fields like **evidence** and more complex value reasoning

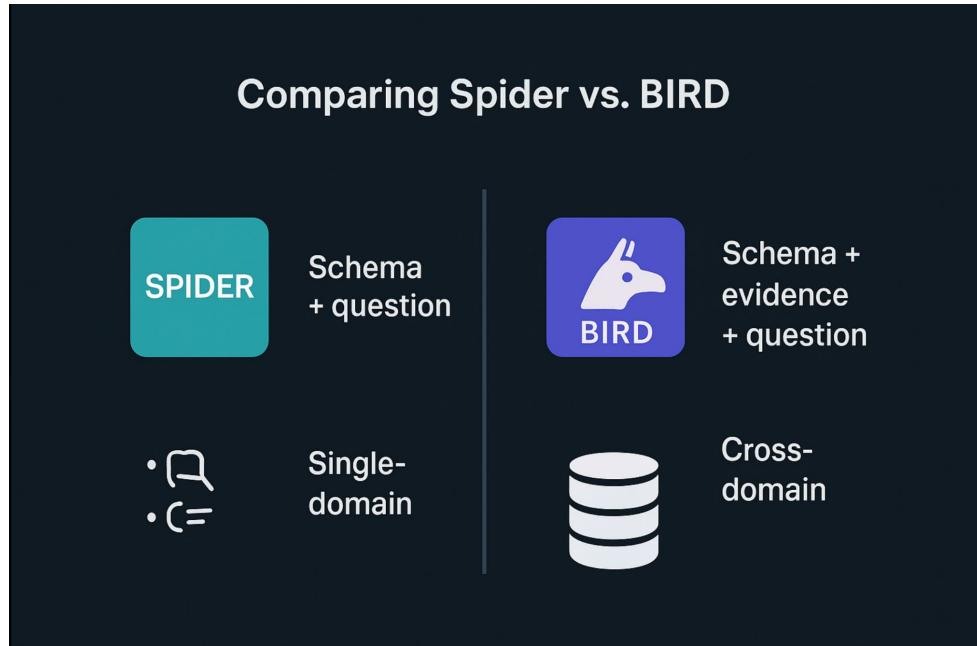
**Pipeline Adjustments**

- New `preprocess_bird.py` to build prompts from `train.json`, `dev.json`, and `/train_databases`, `/dev_databases`.
- Careful schema pruning + delexicalization to fit context window.
- Evaluate with **EX** (Execution Accuracy).



## Phase 3 – Comparing CodeLlama on Spider vs BIRD

- Task Setup:
  - Same base model (**CodeLlama-7B**) and similar fine-tuning strategy.
  - Different data characteristics and evaluation metrics.
- Spider vs BIRD:
  - **Spider**: smaller, cleaner schemas; strong focus on cross-domain generalization; evaluated mainly with **EX**.
  - **BIRD**: larger, real-world databases; more complex value reasoning and evidence; evaluated with **EX**.
- Key Insights:
  - CodeLlama achieves **higher EX on Spider**, showing it handles curated academic benchmarks well.
  - On BIRD, **EX highlights partial correctness** revealing where the model struggles with scale and value reasoning.
  - Together, the two datasets show the trade-off between **benchmark performance** and **real-world readiness**.



# Evaluation Metrics

## Setup

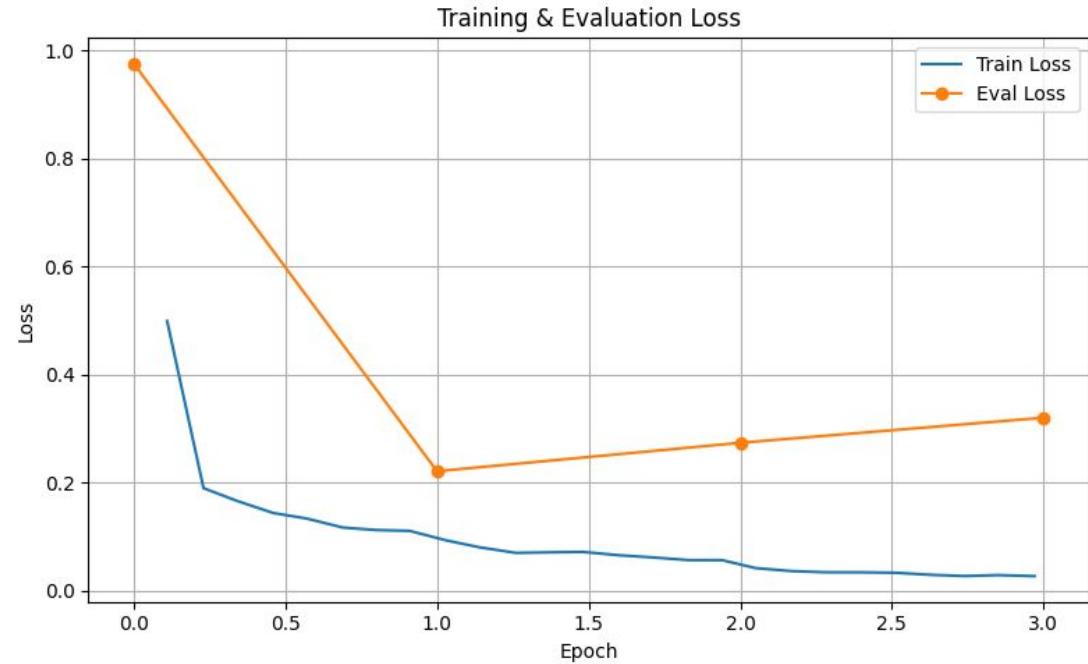
- **Base Model:** IR-NET (Microsoft Research), CodeLlama-7B-hf (Meta)
- **Fine-Tuning:** LoRA adapters
- **Datasets:** Spider 1.0 and BIRD
- **Metrics:** **Execution Accuracy (EX)** → Compares results after executing both queries on real databases.

Metric	Execution Accuracy (EX) Score
Spider 1.0 (With ML)	52.7%
Spider 1.0 (LLM Fine-tuned)	69.3%
Data Segmented Fine-tuned Spider 1.0	68.9%
BIRD (LLM Fine-tuned)	39.60%

## Training & Evaluation Loss

The training loss (blue) drops quickly, showing strong convergence, while evaluation loss (orange) decreases initially then rises slightly indicating mild overfitting after epoch 1.

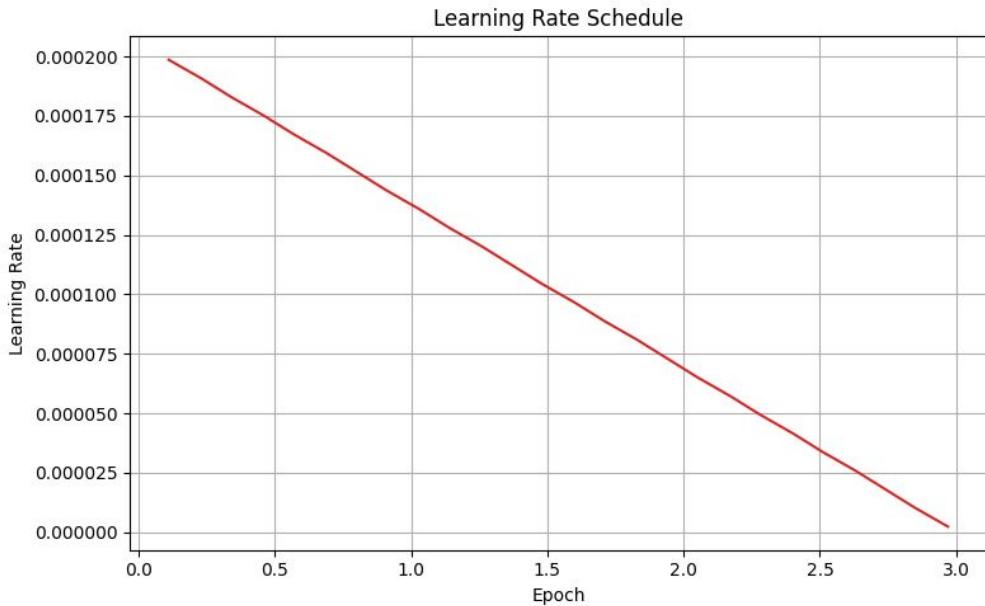
This trend suggests the model learned efficiently, with best generalization likely around epoch 1 and 2.



## Learning Rate Schedule

The learning rate (red) decays linearly from  $2 \times 10^{-4}$  to 0 across 3 epochs, confirming the linear scheduler worked as intended.

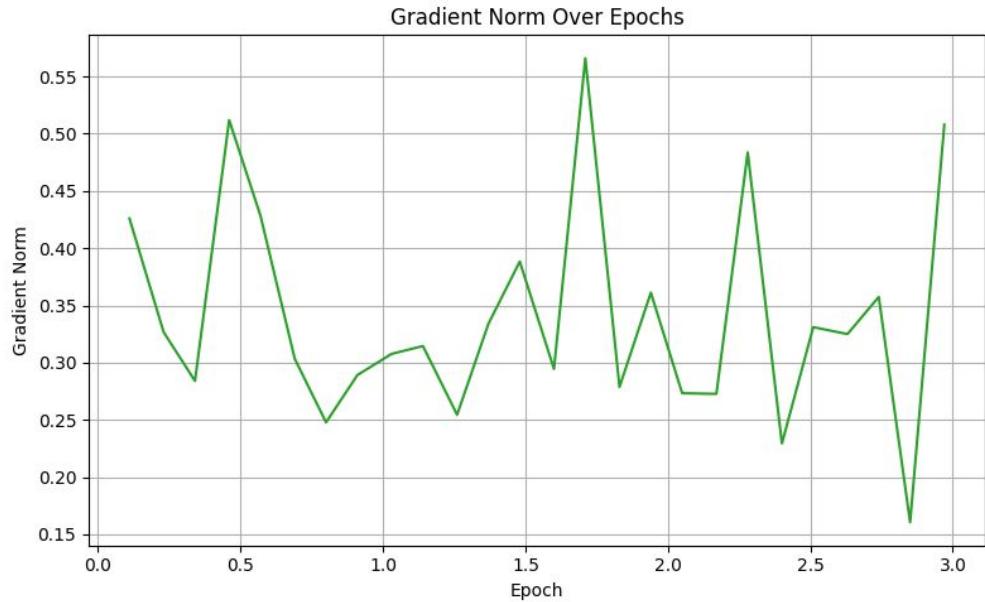
This smooth decay helps stabilize training by allowing large updates early and fine-tuned adjustments later.



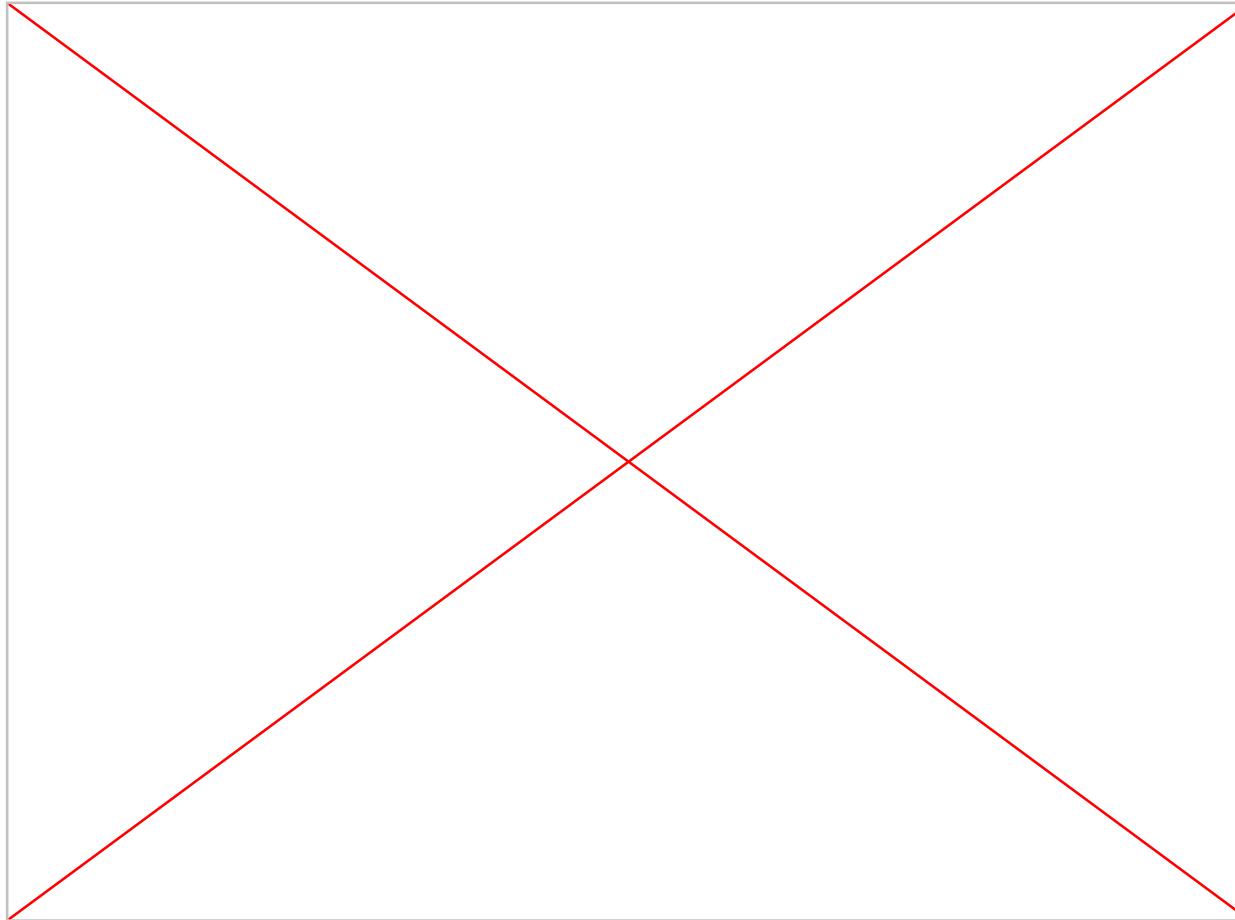
## Gradient Norm Over Epochs

Gradient norms fluctuate moderately ( $\approx 0.2 - 0.55$ ), indicating healthy parameter updates without exploding or vanishing gradients.

This stability shows that our optimizer, LoRA configuration, and learning rate schedule were well-balanced.



# UI Interface Demo



# Future Work and Enhancements

## 1. Hybrid Prompt-plus-Finetune Strategies

- We can explore hybrid approaches that **combine prompt engineering with fine-tuning** to further enhance model performance.
- Through this, we hope to identify a balanced method that maximizes performance while maintaining computational efficiency.

## 2. Experimenting with close-ended LLMs through APIs

- Several LLMs like GPT-4, Deepseek are expected to perform much better given their extremely large number of parameters, however due to their close-ended access, we restricted them for applying to a research project where exploration needs to be significantly conducted.

# Project Plan: Tasks, Deadlines, Division of the Work

Task	Deadline	Responsible
Data Preparation	October 21st	All
Traditional ML Model Training	October 27th	Rebecca Pires dos Santos, Nishtha Chaudhary, Omkar Shelar
LLM Fine Tuning	November 3rd	Jinal Vyas and Vrishir Iyer
Evaluation and Comparison	November 21st	All
Improvement and Research on Advancement Areas	December 5th	All

All contributed in all areas. Responsible refers to the team member responsible for initiating the task and making sure it is completed in the planned time.

# References

- [1] T. Yu, R. Zhang, K. Yang, M. Yasunaga, D. Wang, Z. Li, J. Ma, I. Li, Q. Yao, S. Roman, Z. Zhang, and D. R. Radev, "Spider: A Large-Scale Human-Labeled Dataset for Complex and Cross-Domain Semantic Parsing and Text-to-SQL Task," in Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing (EMNLP), Brussels, Belgium, Oct.–Nov. 2018.
- [2] V. Zhong, C. Xiong, and R. Socher, "Seq2SQL: Generating Structured Queries from Natural Language using Reinforcement Learning," arXiv preprint arXiv:1709.00103, Sep. 2017.
- [3] I. Androultsopoulos, G. D. Ritchie, and P. Thanisch, "Natural Language Interfaces to Databases – An Introduction," arXiv preprint cmp-lg/9503016, Mar. 1995.
- [4] X. Xu, C. Liu, and D. Song, "SQLNet: Generating Structured Queries from Natural Language without Reinforcement Learning," arXiv preprint arXiv:1711.04436, Nov. 2017.
- [5] T. Yu, M. Yasunaga, K. Yang, R. Zhang, D. Wang, Z. Li, and D. R. Radev, "SyntaxSQLNet: Syntax Tree Networks for Complex and Cross-Domain Text-to-SQL Task," arXiv preprint arXiv:1810.05237, Oct. 2018.
- [6] J. Guo, Z. Zhan, Y. Gao, Y. Xiao, J.-G. Lou, T. Liu, and D. Zhang, "Towards Complex Text-to-SQL in Cross-Domain Database with Intermediate Representation," arXiv preprint arXiv:1905.08205, May 2019.
- [7] B. Wang, R. Shin, X. Liu, O. Polozov, and M. Richardson, "RAT-SQL: Relation-Aware Schema Encoding and Linking for Text-to-SQL Parsers," arXiv preprint arXiv:1911.04942, Nov. 2019.
- [8] X. V. Lin, R. Socher, and C. Xiong, "Bridging Textual and Tabular Data for Cross-Domain Text-to-SQL Semantic Parsing," arXiv preprint arXiv:2012.12627, Dec. 2020.
- [9] R. Cao, L. Chen, Z. Chen, Y. Zhao, S. Zhu, and K. Yu, "LGE SQL: Line Graph Enhanced Text-to-SQL Model with Mixed Local and Non-Local Relations," in Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers), Online, Aug. 2021, pp. 2541–2555.
- [10] J. Pennington, R. Socher, and C. D. Manning, "GloVe: Global Vectors for Word Representation," in Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP), Doha, Qatar, Oct. 2014, pp. 1532–1543.
- [11] K. Clark, M.-T. Luong, Q. V. Le, and C. D. Manning, "ELECTRA: Pre-training Text Encoders as Discriminators Rather Than Generators," arXiv preprint arXiv:2003.10555, Mar. 2020.
- [12] T. Yu, R. Zhang, K. Yu, D. Zhang, Z. Wang, W. Zhang, Y. Zhao, D. Lin, M. Sun, D. Zhou, et al., "Spider 2.0: A Large-Scale Cross-Domain Text-to-SQL and Database Grounding Dataset," arXiv preprint arXiv:2209.12391, 2022.
- [13] L. Shi, Z. Tang, N. Zhang, X. Zhang, and Z. Yang, "A Survey on Employing Large Language Models for Text-to-SQL Tasks," ACM Computing Surveys, vol. 58, no. 2, article 54, Sep. 2025, pp. 1–37, doi: 10.1145/3737873. [Online]. Available: <https://doi.org/10.1145/3737873>.
- [14] Jinyang Li, Binyuan Hui, Ge Qu, Jiaxi Yang, Binhu Li, Bowen Li, Bailin Wang, Bowen Qin, Ruiying Geng, Nan Huo, et al. 2024. Can llm already serve as a database interface? a big bench for large-scale database grounded text-to-sqls. Advances in Neural Information Processing Systems 36 (2024) [Online]. Available: <https://arxiv.org/abs/2305.03111>

# Codebase

- [https://github.com/santoreb/Text\\_to\\_SQL](https://github.com/santoreb/Text_to_SQL)

